

frontier shift in one month, for scale
(DeepSWE, 50% success, Apr → May 2026)

◆ 46× (\$4.64 → \$0.10)

sequential search vs best configuration
(DeepSWE, verifier $r \geq 0.8$, equal success)

21.5× (17.1–25.0)

uniform low-cost vs uniform frontier crew
(defect fixes, score difference n.s.)

15.4× [11.4, 19.7]

class-conditional single-shot router vs best configuration
(DeepSWE, 95% of best success)

2.4× (2.0–4.0)

class-aware crew vs best fixed crew
(GitHub issues, score difference 0.00)

1.7× [1.3, 2.6]

1×

2×

5×

10×

20×

50×

100×

cost-reduction factor at equal quality (log scale)